



Red Hat OpenShift AI Self-Managed 2.25

Enabling the model registry component

Enabling the model registry component in Red Hat OpenShift AI Self-Managed

Red Hat OpenShift AI Self-Managed 2.25 Enabling the model registry component

Enabling the model registry component in Red Hat OpenShift AI Self-Managed

Legal Notice

Copyright © 2025 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

Java[®] is a registered trademark of Oracle and/or its affiliates.

XFS[®] is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL[®] is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js[®] is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack[®] Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

After upgrading to OpenShift AI 2.25, ensure that the model registry component is enabled for OpenShift AI administrators and data scientists.

Table of Contents

PREFACE 3

CHAPTER 1. OVERVIEW OF MODEL REGISTRIES AND MODEL CATALOG 4

 1.1. MODEL REGISTRY 4

 1.2. MODEL CATALOG 4

CHAPTER 2. ENABLING THE MODEL REGISTRY COMPONENT 5

PREFACE

As an OpenShift cluster administrator, before data scientists and OpenShift AI administrators in your organization can work with the model registry and model catalog, you must ensure that the model registry component is enabled in OpenShift AI.



NOTE

The model registry component is enabled by default in a new OpenShift AI 2.25 installation. However, if the model registry was not enabled in a previous version of OpenShift AI, you must enable this component after upgrading to OpenShift AI 2.25.

CHAPTER 1. OVERVIEW OF MODEL REGISTRIES AND MODEL CATALOG

A model registry acts as a central repository for administrators and data scientists to register, version, and manage the lifecycle of AI models before configuring them for deployment. A model registry is a key component for AI model governance.

The model catalog provides a curated library where data scientists can discover and evaluate the available generative AI models to find the best fit for their use cases.

1.1. MODEL REGISTRY

A model registry is an important component in the lifecycle of an artificial intelligence/machine learning (AI/ML) model, and is a vital part of any machine learning operations (MLOps) platform or workflow. A model registry acts as a central repository, storing metadata related to machine learning models from development to deployment. This metadata ranges from high-level information like the deployment environment and project, to specific details like training hyperparameters, performance metrics, and deployment events.

A model registry acts as a bridge between model experimentation and serving, offering a secure, collaborative metadata store interface for stakeholders in the ML lifecycle. Model registries provide a structured and organized way to store, share, version, deploy, and track models.

OpenShift AI administrators can create model registries in OpenShift AI and grant model registry access to data scientists. For more information, see [Managing model registries](#).

Data scientists with access to a model registry can use it to store, share, version, deploy, and track models. For more information, see [Working with model registries](#).

1.2. MODEL CATALOG

Data scientists can use the model catalog to discover and evaluate the models that are available and ready for their organization to register, deploy, and customize.

The model catalog provides models from different providers that data scientists can search and discover before they register models in a model registry and deploy them to a model serving runtime. OpenShift AI administrators can configure the available repository sources for models displayed in the model catalog.

OpenShift AI provides a default model catalog, which includes models from providers such as Red Hat, IBM, Meta, Nvidia, Mistral AI, and Google.

For more information about how data scientists can use the model catalog, see [Working with the model catalog](#).

CHAPTER 2. ENABLING THE MODEL REGISTRY COMPONENT

Before data scientists in your organization can work with the model registry and model catalog, you must ensure that the **modelregistry** component is enabled in OpenShift AI.



NOTE

The **modelregistry** component is enabled by default in a new OpenShift AI 2.25 installation. However, if the model registry was not enabled in a previous version of OpenShift AI, you must enable this component after upgrading to OpenShift AI 2.25.

Prerequisites

- You have cluster administrator privileges for your OpenShift cluster.
- You have access to the data science cluster.
- You have installed the Red Hat OpenShift AI Operator on your OpenShift cluster.
- You have sufficient resources. For more information about the minimum resources required to use OpenShift AI, see [Installing and deploying OpenShift AI](#) (for disconnected environments, see [Deploying OpenShift AI in a disconnected environment](#)).

Procedure

1. In the OpenShift console, click **Operators → Installed Operators**.
2. Search for the **Red Hat OpenShift AI Operator** version 2.14+, and then click the Operator name to open the Operator details page.
3. Click the **Data Science Cluster** tab.
4. Click the default instance name (for example, **default-dsc**) to open the instance details page.
5. Click the **YAML** tab to show the instance specifications.
6. Find the **spec.components** section, and then add or update it to include the following **modelregistry** component entry, with the **managementState** field set to **Managed**, and the **registriesNamespace** field set to **rhoai-model-registries**:

```
modelregistry:
  managementState: Managed
  registriesNamespace: rhoai-model-registries
```

7. Click **Save**.

Verification

- Confirm that the model registry namespace was created successfully:
 - a. In the OpenShift console, click **Home → Projects**.
 - b. Confirm that the **rhoai-model-registries** namespace is displayed in the **Projects** drop-down list.

- Check the status of the **model-registry-operator-controller-manager** pod:
 - a. In the OpenShift console, from the **Project** list, select **redhat-ods-applications**.
 - b. Click **Workloads → Deployments**.
 - c. Search for the **model-registry-operator-controller-manager** deployment.
 - d. Check the status:
 - i. Click the deployment name to open the deployment details page.
 - ii. Click the **Pods** tab.
 - iii. View the pod status.
When the status of the **model-registry-operator-controller-manager-*<pod-id>*** pod is **Running**, the pod is ready to use.

Next step

OpenShift AI administrators can create, delete, and manage permissions for model registries. For more information, see [Managing model registries](#).