



Red Hat OpenShift AI Self-Managed 2.25

Working with model registries

Working with model registries in Red Hat OpenShift AI Self-Managed

Legal Notice

Copyright © 2025 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

Java[®] is a registered trademark of Oracle and/or its affiliates.

XFS[®] is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL[®] is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js[®] is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack[®] Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

As a data scientist in OpenShift AI, you can store, share, version, deploy, and track models using the model registry feature.

Table of Contents

PREFACE	3
CHAPTER 1. OVERVIEW OF MODEL REGISTRIES AND MODEL CATALOG	4
1.1. MODEL REGISTRY	4
1.2. MODEL CATALOG	4
CHAPTER 2. WORKING WITH MODEL REGISTRIES	5
2.1. REGISTERING A MODEL FROM THE DASHBOARD	5
2.2. REGISTERING A MODEL FROM THE MODEL CATALOG	6
2.3. REGISTERING A MODEL VERSION	7
2.4. VIEWING REGISTERED MODELS	9
2.5. VIEWING REGISTERED MODEL VERSIONS	9
2.6. EDITING MODEL METADATA IN A MODEL REGISTRY	10
2.7. EDITING MODEL VERSION METADATA IN A MODEL REGISTRY	11
2.8. DEPLOYING A MODEL VERSION FROM A MODEL REGISTRY	12
2.9. DEPLOYING A MODEL FROM THE MODEL CATALOG	13
2.10. EDITING THE DEPLOYMENT PROPERTIES OF A DEPLOYED MODEL VERSION FROM A MODEL REGISTRY	16
2.10.1. Editing the deployment properties of a model deployed by using the multi-model serving platform	16
2.10.2. Editing the deployment properties of a model deployed by using the single-model serving platform	17
2.11. DELETING A DEPLOYED MODEL VERSION FROM A MODEL REGISTRY	19
2.12. ARCHIVING A MODEL	20
2.13. ARCHIVING A MODEL VERSION	20
2.14. RESTORING A MODEL	21
2.15. RESTORING A MODEL VERSION	22

PREFACE

As a data scientist in OpenShift AI, you can store, share, version, deploy, and track models using the model registry feature.

CHAPTER 1. OVERVIEW OF MODEL REGISTRIES AND MODEL CATALOG

A model registry acts as a central repository for administrators and data scientists to register, version, and manage the lifecycle of AI models before configuring them for deployment. A model registry is a key component for AI model governance.

The model catalog provides a curated library where data scientists can discover and evaluate the available generative AI models to find the best fit for their use cases.

1.1. MODEL REGISTRY

A model registry is an important component in the lifecycle of an artificial intelligence/machine learning (AI/ML) model, and is a vital part of any machine learning operations (MLOps) platform or workflow. A model registry acts as a central repository, storing metadata related to machine learning models from development to deployment. This metadata ranges from high-level information like the deployment environment and project, to specific details like training hyperparameters, performance metrics, and deployment events.

A model registry acts as a bridge between model experimentation and serving, offering a secure, collaborative metadata store interface for stakeholders in the ML lifecycle. Model registries provide a structured and organized way to store, share, version, deploy, and track models.

OpenShift AI administrators can create model registries in OpenShift AI and grant model registry access to data scientists. For more information, see [Managing model registries](#).

Data scientists with access to a model registry can use it to store, share, version, deploy, and track models. For more information, see [Working with model registries](#).

1.2. MODEL CATALOG

Data scientists can use the model catalog to discover and evaluate the models that are available and ready for their organization to register, deploy, and customize.

The model catalog provides models from different providers that data scientists can search and discover before they register models in a model registry and deploy them to a model serving runtime. OpenShift AI administrators can configure the available repository sources for models displayed in the model catalog.

OpenShift AI provides a default model catalog, which includes models from providers such as Red Hat, IBM, Meta, Nvidia, Mistral AI, and Google.

For more information about how data scientists can use the model catalog, see [Working with the model catalog](#).

CHAPTER 2. WORKING WITH MODEL REGISTRIES

2.1. REGISTERING A MODEL FROM THE DASHBOARD

As a data scientist, you can register a model from the OpenShift AI dashboard and create the first version of the new model.

Prerequisites

- You are logged in to Red Hat OpenShift AI.
- You have access to an available model registry in your deployment.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that you want to register a model in.
3. Click **Register model**.
The **Register model** dialog opens.
4. In the **Model details** section, configure details to apply to all versions of the model:
 - a. In the **Model name** field, enter a name for the model.
 - b. Optional: In the **Model description** field, enter a description for the model.
5. In the **Version details** section, enter details to apply to the first version of the model:
 - a. In the **Version name** field, enter a name for the model version.
 - b. Optional: In the **Version description** field, enter a description for the first version of the model.
 - c. In the **Source model format** field, enter the name of the model format, for example, **ONNX**.
 - d. In the **Source model format version** field, enter the version of the model format.
6. In the **Model location** section, specify the location of the model by providing either object storage details, or a URI.
 - a. To provide object storage details, ensure that the **Object storage** radio button is selected.
 - i. To autofill the details of an existing connection:
 - A. Click **Autofill from connection**.
 - B. In the **Autofill from connection** dialog that opens, from the **Project** drop-down list, select the data science project that contains the connection.
 - C. From the **Connection name** drop-down list, select the connection that you want to use.
This list contains only object storage types which contain a bucket.

- D. Click **Autofill**.
- ii. Alternatively, manually fill out your object storage details:
 - A. In the **Endpoint** field, enter the endpoint of your S3-compatible object storage bucket.
 - B. In the **Bucket** field, enter the name of your S3-compatible object storage bucket.
 - C. In the **Region** field, enter the region of your S3-compatible object storage account.
 - D. In the **Path** field, enter a path to a model or folder. This path cannot point to a root folder.
- b. To provide a URI, ensure that the **URI** radio button is selected.
 - i. In the **URI** field, enter the URI for the model.



IMPORTANT

Deployment of models that are registered by using a URI is currently supported for public OCI repositories only.

7. Click **Register model**.

Verification

- The new model and version details are displayed on the **Details** tab for the model version.
- The new model and version are displayed on the **Model registry** page.

2.2. REGISTERING A MODEL FROM THE MODEL CATALOG

As a data scientist, you can register models directly from the model catalog and create the first version of the new model.

Prerequisites

- You are logged in to Red Hat OpenShift AI.
- You have access to an available model registry in your deployment.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model catalog**.
2. In the drop-down list, select from the available catalog sources that have been configured by your administrator. The **Default Catalog** is displayed by default.



NOTE

OpenShift cluster administrators can configure additional model catalog sources. For more details, see the [Kubeflow Model Registry community documentation on configuring catalog sources](#).

3. Use the search bar to find a model in the catalog. You can enter text to search by model name, description, or provider.
4. Click the name of a model to view the model details page.
5. Click **Register model**.
6. From the **Model registry** drop-down list, select the model registry that you want to register the model in.
7. In the **Model details** section, configure details to apply to all versions of the model:
 - a. Optional: In the **Model name** field, update the name of the model.
 - b. Optional: In the **Model description** field, update the description of the model.
8. In the **Version details** section, enter details to apply to the first version of the model:
 - a. In the **Version name** field, enter a name for the model version.
 - b. Optional: In the **Version description** field, enter a description for the first version of the model.
 - c. In the **Source model format** field, enter the name of the model format, for example, **ONNX**.
 - d. In the **Source model format version** field, enter the version of the model format.
9. In the **Model location** section, the URI of the model is displayed.
10. Click **Register model**.

Verification

- The new model details and version are displayed on the **Overview** tab on the model details page.
- The new model and version are displayed on the **Model registry** page.

2.3. REGISTERING A MODEL VERSION

You can register a new model version.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- You have access to an available model registry in your deployment.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that you want to register a model version in.

3. In the **Model name** column, click the name of the model that you want to register a new version of.
The details page for the model opens.
4. Click the **Versions** tab, and then click **Register new version**.
5. In the **Version details** section, enter details to apply to this version of the model:
 - a. In the **Version name** field, enter a name for the model version.
 - b. Optional: In the **Version description** field, enter a description for this version of the model.
 - c. In the **Source model format** field, enter the name of the model format, for example, **ONNX**.
 - d. In the **Source model format version** field, enter the version of the model format.
6. In the **Model location** section, specify the location of the model by providing either object storage details, or a URI.
 - a. To provide object storage details, ensure that the **Object storage** radio button is selected.
 - i. To autofill the details of an existing connection:
 - A. Click **Autofill from connection**.
 - B. In the **Autofill from connection** dialog that opens, from the **Project** drop-down list, select the data science project that contains the connection.
 - C. From the **Connection name** drop-down list, select the connection that you want to use.
This list contains only object storage types which contain a bucket.
 - D. Click **Autofill**.
 - ii. Alternatively, manually fill out your object storage details:
 - A. In the **Endpoint** field, enter the endpoint of your S3-compatible object storage bucket.
 - B. In the **Bucket** field, enter the name of your S3-compatible object storage bucket.
 - C. In the **Region** field, enter the region of your S3-compatible object storage account.
 - D. In the **Path** field, enter a path to a model or folder. This path cannot point to a root folder.
 - b. To provide a URI, ensure that the **URI** radio button is selected.
 - i. In the **URI** field, enter the URI for the model.

**IMPORTANT**

Deployment of models that are registered by using a URI is currently supported for public OCI repositories only.

7. Click **Register new version**.

Verification

- The new model version is displayed in the **Latest versions** section on the **Overview** tab on the model details page.
- The new model version is displayed in the **Latest version** column on the **Model registry** page.

2.4. VIEWING REGISTERED MODELS

You can view the details of models registered in OpenShift AI, such as registered versions, deployments, and metadata associated with the model.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least 1 registered model.
- You have access to the model registry that contains the model that you want to view.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that contains the model that you want to view.
3. The **Model registry** page provides a high-level view of registered models, including the model name, latest version, deployments, labels, last modified timestamp, and owner of each model. Models are sorted by their **Last modified** timestamp by default.
4. Use the search bar to find a model in the list. You can filter with a keyword by default by entering a model name, description, or label. Alternatively, click the search bar drop-down list and select **Owner** to filter by entering a model owner.
Searching by keyword performs a search across the name, description, and labels of registered models and their versions.
5. Click the name of a model to view the details page for the model:
 - a. On the **Overview** tab, you can view model metadata such as labels, description, owner, model ID, last modified and created timestamps, and custom properties, along with latest versions and deployments.
 - b. On the **Versions** tab, you can view the registered versions of the model.
 - c. On the **Deployments** tab, you can view deployments initiated from the model registry for this model.

Verification

- You can view information about the selected model on the details page for the model.

2.5. VIEWING REGISTERED MODEL VERSIONS

You can view the details of model versions that are registered in OpenShift AI, such as the version metadata and deployment information.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least 1 registered model.
- You have access to the model registry that contains the model version that you want to view.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that contains the model version that you want to view.
3. Click the name of a model to view **Overview** tab on the model details page, which includes the latest model versions and deployments.
4. On the **Versions** tab, you can view the registered versions of the model. Versions are sorted by their **Last modified** timestamp by default.
5. Use the search bar to find a version in the list. You can filter with a keyword by default by entering a model name, description, or label. Alternatively, click the search bar drop-down list and select **Author** to filter by entering a model owner. Searching by keyword performs a search across the name, description, and labels of registered models and their versions.
6. Click the name of a version to view the version details page.
7. On the **Details** tab, you can view the **Version details** metadata, such as labels, description, custom properties, version ID, author, and last modified and registered timestamps. This also includes where the model is registered from, model location, and model format information. You can also click **Model details** to view non-version metadata, such as labels, description, owner, model ID, last modified and created timestamps, and custom properties.
8. On the **Deployments** tab, you can view deployments initiated from the model registry for this version.
 - a. Click the name of a deployment to open its metrics page.
For information about model metrics on the single-model serving platform, see [Viewing performance metrics for a deployed model](#). For information about model metrics on the multi-model serving platform, see [Viewing model-serving runtime metrics for the multi-model serving platform](#).

Verification

- You can view the details of registered model versions on the **Model registry** page.

2.6. EDITING MODEL METADATA IN A MODEL REGISTRY

You can edit the metadata of models registered in OpenShift AI, such as the model description, labels, and custom properties. Editing model metadata affects all versions of the model.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least 1 registered model.
- You have access to the model registry that contains the model that you want to edit.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that contains the model that you want to edit.
3. The **Model registry** page provides a high-level view of registered models, including the model name, latest version, deployments, labels, last modified timestamp, and owner of each model.
4. Click the name of a model to view the model details page.
5. On the **Overview** tab, you can edit metadata for the model.
 - a. In the **Labels** section, click **Edit** to edit the labels of the model, for example, **text-to-text**.
 - b. In the **Description** section, click **Edit** to edit the description of the model.
 - c. In the **Properties** section, click **Add property** to add a new property to the model, for example, **Key: license, Value: apache**.

TIP

If you enter any property value as a URL, this is displayed as a clickable link in the **Properties** section, for example: <https://www.apache.org/licenses/LICENSE-2.0>.

- i. To edit an existing property, click the action menu (⋮) beside the property, and then click **Edit**.
- ii. To delete a property, click the action menu (⋮) beside the property, and then click **Delete**.

Verification

- You can view the updated metadata on the details page for the model.

2.7. EDITING MODEL VERSION METADATA IN A MODEL REGISTRY

You can edit the metadata of model versions that are registered in OpenShift AI, such as the version's description, labels, and custom properties. Editing model version metadata affects that model version only.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least 1 registered model.

- You have access to the model registry that contains the model version that you want to edit.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that contains the model version that you want to edit.
3. Click the name of a model to view the model details page.
4. Click the **Versions** tab to view the available model versions.
5. Click a version name to view the version details page.
6. In the **Version details** section, you can edit the version metadata.
 - a. In the **Labels** section, click **Edit** to edit the labels of the version, for example, **text-to-text**.
 - b. In the **Description** section, click **Edit** to edit the description of the version.
 - c. In the **Properties** section, click **Add property** to add a new property to the version, for example, **Key: license, Value: apache**.

TIP

If you enter any property value as a URL, this is displayed as a clickable link in the **Properties** section, for example: <https://www.apache.org/licenses/LICENSE-2.0>.

- i. To edit an existing property, click the action menu (⋮) beside the property, and then click **Edit**.
 - ii. To delete a property, click the action menu (⋮) beside the property, and then click **Delete**.
- d. In the **Model format** section, click **Edit** to edit the format of the model version, for example, **ONNX**.
 - e. In the **Model format version** section, click **Edit** to edit the format version of the model version.

Verification

- You can view the updated metadata on the details page for the model version.

2.8. DEPLOYING A MODEL VERSION FROM A MODEL REGISTRY

You can deploy a version of a registered model directly from a model registry.

Prerequisites

- An available model registry exists in your deployment, and contains at least 1 registered model.
- To deploy a model version by using the single-model serving platform, you have fulfilled the prerequisites described in [Deploying models on the single-model serving platform](#).

- To deploy a model version by using the multi-model serving platform, you have fulfilled the prerequisites described in [Deploying a model by using the multi-model serving platform](#).

Procedure

1. From the OpenShift AI dashboard, click **Models** → **Model registry**.
2. From the **Model registry** drop-down list, select the model registry from which you want to deploy a model version.
3. In the **Model name** column, click the name of the model that contains the version that you want to deploy.
The details page for the model version opens.
4. Click the action menu () beside the model version that you want to deploy.
5. Click **Deploy**.
6. In the **Deploy model** dialog, configure properties for deploying the model.
 - a. From the **Project** drop-down list, select a target project.
 - b. Click **Deploy**.
7. Configure the following properties for deploying your model:
 - a. From the **Project** drop-down list, select a project in which to deploy your model.
 - b. Optional: In the **Model deployment name** field, enter a unique name for your model deployment. This field is autofilled with a value that contains the model name by default. This will be the name of the inference service that is created when the model is deployed.
8. Configure the remaining properties for deploying your model, as described in [Deploying a model by using the multi-model serving platform](#) or [Deploying models on the single-model serving platform](#).
 - a. Click **Deploy**.

Verification

- The model deployment is displayed on the **Models** → **Model Deployments** page.
- The model deployment is displayed in the **Latest deployments** section of the model details page.
- The model version is displayed on the **Deployments** tab for the model.
- You can edit the model version deployment by clicking the action menu () beside it, and then clicking **Edit**.
- You can delete the model version deployment by clicking the action menu () beside it, and then clicking **Delete**.

2.9. DEPLOYING A MODEL FROM THE MODEL CATALOG

You can deploy models directly from the model catalog.

**NOTE**

OpenShift AI model serving deployments use the global cluster pull secret to pull models in ModelCar format from the catalog.

For more information about using pull secrets in OpenShift, see [Updating the global cluster pull secret](#) in the OpenShift documentation.

Prerequisites

- You have completed the prerequisites in [Deploying models on the single-model serving platform](#).
- The model registry component is enabled in your OpenShift AI deployment. For more information, see [Enabling the model registry component](#).

Procedure

1. From the OpenShift AI dashboard, click **Models → Model catalog**.
2. In the drop-down list, select from the available catalog sources that have been configured by your administrator. The **Default Catalog** is displayed by default.

**NOTE**

OpenShift cluster administrators can configure additional model catalog sources. For more details, see the KubeFlow Model Registry community documentation on [configuring catalog sources](#).

3. Use the search bar to find a model in the catalog. You can enter text to search by model name, description, or provider.
4. Click the name of a model to view the model details page.
5. Click **Deploy model** to display the **Deploy model** dialog.
6. From the **Project** drop-down list, select a project in which to deploy your model.

**NOTE**

Models using OCI storage can only be deployed on the single-model serving platform. Projects using the multi-model serving platform do not appear in the project list.

7. In the **Model deployment** section:
 - a. Optional: In the **Model deployment name** field, enter a unique name for your model deployment. This field is autofilled with a value that contains the model name by default. This is the name of the inference service created when the model is deployed.
 - b. Optional: Click **Edit resource name**, and then enter a specific resource name for the model deployment in the **Resource name** field. By default, the resource name matches the name of the model deployment.



IMPORTANT

Resource names are what your resources are labeled as in OpenShift. Your resource name cannot exceed 253 characters, must consist of lowercase alphanumeric characters or -, and must start and end with an alphanumeric character. Resource names are not editable after creation.

The resource name must not match the name of any other model deployment resource in your OpenShift cluster.

- c. From the **Serving runtime** list, select a model-serving runtime that is installed and enabled in your OpenShift AI deployment. If project-scoped runtimes exist, the **Serving runtime** list includes subheadings to distinguish between global runtimes and project-scoped runtimes.
- d. From the **Model framework** list, select a framework for your model.



NOTE

The **Model framework** list shows only the frameworks that are supported by the model-serving runtime that you specified when you deployed your model.

- 8. From the **Deployment mode** list, select **KServe RawDeployment** or **Knative Serverless**. For more information about deployment modes, see [About KServe deployment modes](#).
 - a. In the **Number of model server replicas to deploy** field, specify a value.
 - b. From the **Model server size** list, select a value.
 - c. If you have created a hardware profile, select a hardware profile from the **Hardware profile** list. If project-scoped hardware profiles exist, the **Hardware profile** list includes subheadings to distinguish between global hardware profiles and project-scoped hardware profiles.



IMPORTANT

By default, hardware profiles are hidden from appearing in the dashboard navigation menu and user interface. In addition, user interface components associated with the deprecated accelerator profiles functionality are still displayed. To show the **Settings → Hardware profiles** option in the dashboard navigation menu and the user interface components associated with hardware profiles, set the **disableHardwareProfiles** value to **false** in the **OdhDashboardConfig** custom resource (CR) in OpenShift. For more information about setting dashboard configuration options, see [Customizing the dashboard](#).

- d. In the **Model route** section, select the **Make deployed models available through an external route** checkbox to make your deployed models available to external clients.
- e. In the **Token authentication** section, select the **Require token authentication** checkbox to require token authentication for your model server. To finish configuring token authentication, perform the following actions:
 - i. In the **Service account name** field, enter a service account name for which the token will be generated. The generated token is created and displayed in the **Token secret** field when the model server is configured.

- ii. To add an additional service account, click **Add a service account** and enter another service account name.
9. In the **Source model location** section, select **Current URI** to deploy the selected model from the catalog.
10. Optional: Customize the runtime parameters in the **Configuration parameters** section:
 - a. Modify the values in **Additional serving runtime arguments** to define how the deployed model behaves.
 - b. Modify the values in **Additional environment variables** to define variables in the model's environment.
11. Click **Deploy**.

Verification

- The model deployment is displayed on the **Models → Model Deployments** page.
- The model deployment is displayed in the **Latest deployments** section of the model details page.
- The model deployment is displayed on the **Deployments** tab for the model version.

Additional resources

- For more information about deployment options on the single-model serving platform, see [Deploying models on the single-model serving platform](#).

2.10. EDITING THE DEPLOYMENT PROPERTIES OF A DEPLOYED MODEL VERSION FROM A MODEL REGISTRY

You can edit model version deployment properties from a model registry for models that were deployed from the registry. For example, you can change the deployment name, model framework, and source model location details.

2.10.1. Editing the deployment properties of a model deployed by using the multi-model serving platform

You can edit the deployment properties of a deployed model version from a model registry. For example, you can change the deployment name, model framework, and source model location details.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least 1 registered and deployed model version.
- You have access to the model registry that contains the model version deployment that you want to edit.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that contains the model deployment that you want to edit.
3. In the **Model name** column, click the name of the model that contains the deployment that you want to edit.
The details page for the model opens.
4. Click the name of the model version with the deployment that you want to edit.
5. Click **Deployments**
6. Click the action menu () beside the model deployment that you want to edit.
7. Click **Edit**.
8. In the **Edit model** dialog, edit the model deployment properties:
 - a. In the **Model deployment name** field, enter a new, unique name for your model deployment.
 - b. From the **Model framework** list, select a different framework for your model.

**NOTE**

The **Model framework** list shows only the frameworks that are supported by the model serving runtime that you specified when you configured your model server.

- c. Edit the connection by specifying an existing connection, or by creating a new connection.
- d. Click **Redeploy**.

Verification

- The model redeloys and is displayed with updated details on the **Deployments** tab for the model version.

2.10.2. Editing the deployment properties of a model deployed by using the single-model serving platform

You can edit the deployment properties of a deployed model version from a model registry. For example, you can change the deployment name, model framework, number of model server replicas, model server size, and source model location details.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least 1 registered and deployed model version.
- You have access to the model registry that contains the model version deployment that you want to edit.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that contains the model deployment that you want to edit.
3. In the **Model name** column, click the name of the model that contains the deployment that you want to edit.
The details page for the model opens.
4. Click the name of the model version with the deployment that you want to edit.
5. Click **Deployments**
6. Click the action menu () beside the model deployment that you want to edit.
7. Click **Edit**.
8. In the **Edit model** dialog, edit the model deployment properties:
 - a. In the **Model deployment name** field, enter a new, unique name for your model deployment.
 - b. From the **Model framework** list, select a different framework for your model.



NOTE

The **Model framework** list shows only the frameworks that are supported by the model serving runtime that you specified when you deployed your model.

- c. In the **Number of model server replicas to deploy** field, specify a value.
- d. From the **Model server size** list, select a value.
- e. In the **Model route** section, select the **Make deployed models available through an external route** checkbox to make your deployed models available to external clients.
- f. In the **Token authentication** section, select the **Require token authentication** checkbox to require token authentication for your model server. To finish configuring token authentication, perform the following actions:
 - i. In the **Service account name** field, enter a service account name for which the token will be generated. The generated token is created and displayed in the **Token secret** field when the model server is configured.
 - ii. To add an additional service account, click **Add a service account** and enter another service account name.
- g. Edit the connection by specifying an existing connection, or by creating a new connection.
- h. Customize the runtime parameters in the **Configuration parameters** section:
 - i. Modify the values in **Additional serving runtime arguments** to define how the deployed model behaves.
 - ii. Modify the values in **Additional environment variables** to define variables in the

- ii. Modify the values in **Additional environment variables** to define variables in the model's environment.

The **Configuration parameters** section shows predefined serving runtime parameters, if any are available.



NOTE

Do not modify the port or model serving runtime arguments, because they require specific values to be set. Overwriting these parameters can cause the deployment to fail.

- i. Click **Redeploy**.

Verification

- The model redeployes and is displayed with updated details on the **Deployments** tab for the model version.

2.11. DELETING A DEPLOYED MODEL VERSION FROM A MODEL REGISTRY

You can delete the deployments of model versions from a model registry.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least 1 registered model with a deployed model version.
- You have access to the model registry that contains the model version deployment that you want to delete.

Procedure


1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that contains the deployment that you want to delete.
3. Click the name of a model to view more details.
The details page for the model opens.
4. Click the name of the model version with the deployment that you want to delete.
The details page for the model version opens.
5. Click **Deployments**.
6. To delete a deployment, click the action menu (⋮) beside the deployment, and then click **Delete**.
The **Delete deployed model?** dialog opens.
7. Enter the name of the model deployment in the text field to confirm that you intend to delete it.
8. Click **Delete deployed model**.

Verification

- The model deployment is no longer displayed on the **Deployments** tab for the model version.

2.12. ARCHIVING A MODEL

You can archive a model that you no longer require. The model and all of its versions will be archived and unavailable for use unless it is restored.



IMPORTANT

Models with deployed versions cannot be archived. To archive a model, you must first delete all deployments of its registered versions from the **Models → Model deployments** page.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least 1 registered model.
- You have access to the model registry that contains the model that you want to archive.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that you want to archive a model in.
3. Click the action menu () beside the model that you want to archive.
4. Click **Archive model**.
5. In the **Archive model?** dialog that is displayed, enter the name of the model in the text field to confirm that you intend to archive it.
6. Click **Archive**.

Verification

- The model is no longer visible on the **Model registry** page.
- The model is displayed on the archived models page for the model registry.

2.13. ARCHIVING A MODEL VERSION

You can archive a model version that you no longer require. The model version will be archived and unavailable for use unless it is restored.



IMPORTANT

Deployed model versions cannot be archived. To archive a model version, you must first delete all deployments of the version from the **Models → Model deployments** page.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least 1 registered model.
- You have access to the model registry that contains the model version that you want to archive.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that you want to archive a model in.
3. In the **Model name** column, click the name of the model that contains the version that you want to archive.
The details page for the model version opens.
4. Click the action menu (⋮) beside the version that you want to archive.
5. Click **Archive model version**.
6. In the **Archive version?** dialog that opens, enter the name of the model version in the text field to confirm that you intend to archive it.
7. Click **Archive**.

Verification

- The model version is no longer visible on the details page for the model.
- The model version is displayed on the archived versions page for the model.

2.14. RESTORING A MODEL

You can restore an archived model. The model and all of its versions will be restored and returned to the registered models list.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least one archived model.
- You have access to the model registry that contains the model that you want to restore.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that you want to restore a model in.
3. Click the action menu (⋮) beside the **Register model** button, and then click **View archived models**.

The archived models page for the model registry opens.

4. Click the action menu (⋮) beside the model that you want to restore.
5. Click **Restore model**.
6. In the **Restore model?** dialog that is displayed, click **Restore**.

Verification

- The model is displayed on the **Model registry** page.
- The model is no longer displayed on the archived models page for the model registry.

2.15. RESTORING A MODEL VERSION

You can restore an archived model version. The model version will be restored and returned to the versions list for the model.

Prerequisites

- You have logged in to Red Hat OpenShift AI.
- An available model registry exists in your deployment, and contains at least one archived model version.
- You have access to the model registry that contains the model version that you want to restore.

Procedure

1. From the OpenShift AI dashboard, click **Models → Model registry**.
2. From the **Model registry** drop-down list, select the model registry that you want to restore a model version in.
3. In the **Model name** column, click the name of the model that contains the version that you want to restore.
The details page for the model version opens.
4. Click the action menu (⋮) beside the **Register new version** button, and then click **View archived versions**.
The archived versions page for the model opens.
5. Click the action menu (⋮) beside the version that you want to restore.
6. Click **Restore version**.
7. In the **Restore version?** dialog that opens, click **Restore**.
The details page for the version opens.

Verification

- The model version is displayed on the details page for the model.
- The model is no longer displayed on the archived versions page for the model.

